

# High Performance Numerical Linear Algebra

Victor Eijkhout

- Dense matrix-vector product
- Sparse matrix-vector product
- Nested dissection
- Parallel preconditioners

# Dense matrix-vector product

# Parallel matrix-vector product; dense

- Assume a division by block rows
- Every processor  $p$  has a set of row indices  $I_p$

Mvp on processor  $p$ :


$$\forall_i: y_i = \sum_j a_{ij}x_j$$

$$\forall_i: y_i = \sum_q \sum_{j \in I_q} a_{ij}x_j$$

Local and remote parts:

$$\forall_i: y_i = \sum_{j \in I_p} a_{ij} x_j + \sum_{q \neq p} \sum_{j \in I_q} a_{ij} x_j$$

Local part  $I_p$  can be executed right away,  $I_q$  requires communication.



Combine:

**Input:** Processor number  $p$ ; the elements  $x_i$  with  $i \in I_p$ ; matrix elements  $A_{ij}$  with  $i \in I_p$ .

**Output:** The elements  $y_i$  with  $i \in I_p$

**for**  $q \neq p$  **do**

    | Send elements of  $x$  from processor  $q$  to  $p$ , receive in buffer  $B_{pq}$ .

**end**

$y_{local} \leftarrow Ax_{local}$

**for**  $q \neq p$  **do**

    |  $y_{local} \leftarrow y_{local} + A_{pq}B_q$

**end**

**Procedure** Parallel MVP( $A, x_{local}, y_{local}, p$ )

Note possible overlap communication and computation

# Cost computation 1.

Algorithm:

Step	Cost (lower bound)
Allgather $x_i$ so that $x$ is available on all nodes Locally compute $y_i = A_i x$	$\approx 2 \frac{n^2}{P} \gamma$

# Allgather

Assume that data arrives over a binary tree:

- latency  $\alpha \log_2 P$
- transmission time, receiving  $n/P$  elements from  $P - 1$  processors



Algorithm with cost:

Step	Cost (lower bound)
Allgather $x_i$ so that $x$ is available on all nodes	$\lceil \log_2(P) \rceil \alpha + \frac{P-1}{P} n \beta \approx \log_2(P) \alpha + n \beta$
Locally compute $y_i = A_i x$	$\approx 2 \frac{n^2}{P} \gamma$

# Parallel efficiency

$$E_p^{1\text{D-row}}(n) = \frac{S_p^{1\text{D-row}}(n)}{p} = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}}.$$

Strong scaling, weak scaling?

# Two-dimensional partitioning

$x_0$ $a_{00}$ $a_{01}$ $a_{02}$ $y_0$ $a_{10}$ $a_{11}$ $a_{12}$ $a_{20}$ $a_{21}$ $a_{22}$ $a_{30}$ $a_{31}$ $a_{32}$	$x_3$ $a_{03}$ $a_{04}$ $a_{05}$ $a_{13}$ $a_{14}$ $a_{15}$ $y_1$ $a_{23}$ $a_{24}$ $a_{25}$ $a_{33}$ $a_{34}$ $a_{35}$	$x_6$ $a_{06}$ $a_{07}$ $a_{08}$ $a_{16}$ $a_{17}$ $a_{18}$ $a_{26}$ $a_{27}$ $a_{28}$ $y_2$ $a_{36}$ $a_{37}$ $a_{38}$	$x_9$ $a_{09}$ $a_{10}$ $a_{19}$ $a_{20}$ $a_{29}$ $a_{30}$ $a_{39}$ $a_{40}$
$x_1$ $a_{40}$ $a_{41}$ $a_{42}$ $y_4$ $a_{50}$ $a_{51}$ $a_{52}$ $a_{60}$ $a_{61}$ $a_{62}$ $a_{70}$ $a_{71}$ $a_{72}$	$x_4$ $a_{43}$ $a_{44}$ $a_{45}$ $a_{53}$ $a_{54}$ $a_{55}$ $y_5$ $a_{63}$ $a_{64}$ $a_{65}$ $a_{73}$ $a_{74}$ $a_{75}$	$x_7$ $a_{46}$ $a_{47}$ $a_{48}$ $a_{56}$ $a_{57}$ $a_{58}$ $a_{66}$ $a_{67}$ $a_{68}$ $y_6$ $a_{76}$ $a_{77}$ $a_{78}$	$x_{10}$ $a_{49}$ $a_{50}$ $a_{59}$ $a_{60}$ $a_{69}$ $a_{70}$ $a_{79}$ $a_{80}$
$x_2$ $a_{80}$ $a_{81}$ $a_{82}$ $y_8$ $a_{90}$ $a_{91}$ $a_{92}$ $a_{10,0}$ $a_{10,1}$ $a_{10,2}$ $a_{11,0}$ $a_{11,1}$ $a_{11,2}$	$x_5$ $a_{83}$ $a_{84}$ $a_{85}$ $a_{93}$ $a_{94}$ $a_{95}$ $y_9$ $a_{10,3}$ $a_{10,4}$ $a_{10,5}$ $a_{11,3}$ $a_{11,4}$ $a_{11,5}$	$x_8$ $a_{86}$ $a_{87}$ $a_{88}$ $a_{96}$ $a_{97}$ $a_{98}$ $a_{10,6}$ $a_{10,7}$ $a_{10,8}$ $y_{10}$ $a_{11,6}$ $a_{11,7}$ $a_{11,8}$	$x_{11}$ $a_{89}$ $a_{90}$ $a_{99}$ $a_{100}$ $a_{10,9}$ $a_{10,10}$ $a_{11,9}$ $a_{11,10}$

# Algorithm

- Collecting  $x_j$  on each processor  $p_{ij}$  by an *allgather* inside the processor columns.
- Each processor  $p_{ij}$  then computes  $y_{ij} = A_{ij}x_j$ .
- Gathering together the pieces  $y_{ij}$  in each processor row to form  $y_i$ , distribute this over the processor row: combine to form a *reduce-scatter*.
- Setup for the next  $A$  or  $A^t$  product

# Analysis 1.

Step	Cost (lower bound)
Allgather $x_i$ 's within columns	$\lceil \log_2(r) \rceil \alpha + \frac{r-1}{p} n \beta \approx \log_2(r) \alpha + \frac{n}{c} \beta$
Perform local matrix-vector multiply	$\approx 2 \frac{n^2}{p} \gamma$
Reduce-scatter $y_i$ 's within rows	

# Reduce-scatter

	$t = 1$	$t = 2$	$t = 3$
$p_0$	$x_0^{(0)}, x_1^{(0)}, x_2^{(0)} \downarrow, x_3^{(0)} \downarrow$	$x_0^{(0:2:2)}, x_1^{(0:2:2)} \downarrow$	$x_0^{(0:3)}$
$p_1$	$x_0^{(1)}, x_1^{(1)}, x_2^{(1)} \downarrow, x_3^{(1)} \downarrow$	$x_0^{(1:3:2)} \uparrow, x_1^{(1:3:2)}$	$x_1^{(0:3)}$
$p_2$	$x_0^{(2)} \uparrow, x_1^{(2)} \uparrow, x_2^{(2)}, x_3^{(2)}$	$x_2^{(0:2:2)}, x_3^{(0:2:2)} \downarrow$	$x_2^{(0:3)}$
$p_3$	$x_0^{(3)} \uparrow, x_1^{(3)} \uparrow, x_2^{(3)}, x_3^{(3)}$	$x_0^{(1:3:2)} \uparrow, x_1^{(1:3:2)}$	$x_3^{(0:3)}$

Time:

$$\lceil \log_2 p \rceil \alpha + \frac{p-1}{p} n(\beta + \gamma).$$

Step	Cost (lower bound)
Allgather $x_i$ 's within columns	$\lceil \log_2(r) \rceil \alpha + \frac{r-1}{p} n \beta \approx \log_2(r) \alpha + \frac{n}{c} \beta$
Perform local matrix-vector multiply	$\approx 2 \frac{n^2}{p} \gamma$
Reduce-scatter $y_i$ 's within rows	$\lceil \log_2(c) \rceil \alpha + \frac{c-1}{p} n \beta + \frac{c-1}{p} n \gamma \approx \log_2(r) \alpha + \frac{n}{c} \beta + \frac{n}{c} \gamma$

# Efficiency

$$E_p^{\sqrt{p} \times \sqrt{p}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{(2\beta + \gamma)}{\gamma}}$$

Weak scaling:

for  $p \rightarrow \infty$  this is  $\approx 1 / \log_2 P$ :

only slowly decreasing.




# LU factorization scaling

Needs a cyclic distribution


# Sparse matrix-vector product

# Sparse matrix-vector product

Difference stencil



induces ghost region:



Limited number of neighbours, limited buffer space

# Scaling

Separately 1D and 2D partitioning of the domain.

# Nested dissection

# Fill-in during LU

Fill-in: index  $(i, j)$  where  $a_{ij} = 0$  but  $\ell_{ij} \neq 0$  or  $u_{ij} \neq 0$ .

2D BVP:  $\Omega$  is  $n \times n$ , gives matrix of size  $N = n^2$ , with bandwidth  $n$ .

Matrix storage  $O(N)$

LU storage  $O(N^{3/2})$

LU factorization work  $O(N^2)$

Cute fact: storage can be computed linear in #nonzeros

# Fill-in is a function of ordering


$$\begin{pmatrix} * & * & \dots & * \\ * & * & & \emptyset \\ \vdots & & \ddots & \\ * & \emptyset & & * \end{pmatrix}$$

After factorization the matrix is dense.

Can this be permuted?



# Domain decomposition



$$\left( \begin{array}{ccccc|ccccc|c}
 \star & \star & & & & & & & & & 0 \\
 \star & \star & \star & & & & & & & & \vdots \\
 & \ddots & \ddots & \ddots & & & & & & & \vdots \\
 & & \star & \star & \star & \emptyset & & & & & 0 \\
 & & & \star & \star & & & & & & \star \\
 \hline
 & & & & & \star & \star & & & & 0 \\
 & & & & & \star & \star & \star & & & \vdots \\
 & & \emptyset & & & \ddots & \ddots & \ddots & & & \vdots \\
 & & & & & & \star & \star & \star & & 0 \\
 & & & & & & & \star & \star & & \star \\
 \hline
 0 & \dots & \dots & 0 & \star & 0 & \dots & \dots & 0 & \star & \star
 \end{array} \right)
 \left. \begin{array}{l} \\ \\ \\ \\ \\ \\ \\ \\ \\ \end{array} \right\} \begin{array}{l} (n^2 - n)/2 \\ (n^2 - n)/2 \\ n \end{array}$$


# DD factorization

$$A^{\text{DD}} = \begin{pmatrix} A_{11} & \emptyset & A_{13} \\ \emptyset & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix} = \begin{pmatrix} I & & \\ & I & \\ A_{31}A_{11}^{-1} & A_{32}A_{22}^{-1} & I \end{pmatrix} \begin{pmatrix} A_{11} & \emptyset & A_{13} \\ & A_{22} & A_{23} \\ & & S \end{pmatrix}$$
$$S = A_{33} - A_{31}A_{11}^{-1}A_{13} - A_{32}A_{22}^{-1}A_{23}$$


Parallelism...

# Graph theory of sparse elimination

$$a_{ij} \leftarrow a_{ij} - a_{ik} a_{kk}^{-1} a_{kj}$$




# Graph theory of sparse elimination



# Graph theory of sparse elimination

$$a_{ij} \leftarrow a_{ij} - a_{ik} a_{kk}^{-1} a_{kj}$$



So inductively  $S$  is dense

# Recursive bisection





Figure : A four-way domain decomposition

$$A^{\text{DD}} = \begin{pmatrix} A_{11} & & & & A_{15} & & A_{17} \\ & A_{22} & & & A_{25} & & A_{27} \\ & & A_{33} & & & A_{36} & A_{37} \\ & & & A_{44} & & A_{46} & A_{47} \\ A_{51} & A_{52} & & & A_{55} & & A_{57} \\ & & A_{63} & A_{64} & & A_{66} & A_{67} \\ A_{71} & A_{72} & A_{73} & A_{74} & A_{75} & A_{76} & A_{77} \end{pmatrix}$$

The domain/operator/graph view is more insightful, don't you think?



# How does this look in reality?



# Complexity

With  $n = \sqrt{N}$ :

- one dense matrix on a separator of size  $n$ , plus
- two dense matrices on separators of size  $n/2$
- $\rightarrow 3/2 n^2$  space and  $5/12 n^3$  time
- and then four times the above with  $n \rightarrow n/2$

$$\begin{aligned}\text{space} &= 3/2 n^2 + 4 \cdot 3/2 (n/2)^2 + \dots \\ &= N(3/2 + 3/2 + \dots) \quad \log n \text{ terms} \\ &= O(N \log N)\end{aligned}$$

$$\begin{aligned}\text{time} &= 5/12 n^3/3 + 4 \cdot 5/12 (n/2)^3/3 + \dots \\ &= 5/12 N^{3/2} (1 + 1/4 + 1/16 + \dots) \\ &= O(N^{3/2})\end{aligned}$$

# More direct factorizations

Minimum degree, multifrontal, . . .

Finding good separators and domain decompositions is tough in general.

# Parallel preconditioners

# Sparse operations in parallel: mvp

Mvp  $y = Ax$

```
for i=1..n  
  y[i] = sum over j=1..n a[i,j]*x[j]
```

In parallel:

```
for i=myfirstrow..mylastrow  
  y[i] = sum over j=1..n a[i,j]*x[j]
```

# How about ILU solve?

Consider  $Lx = y$

```
for i=1..n
  x[i] = (y[i] - sum over j=1..i-1 ell[i,j]*x[j])
        / a[i,i]
```

Parallel code:

```
for i=myfirstrow..mylastrow
  x[i] = (y[i] - sum over j=1..i-1 ell[i,j]*x[j])
        / a[i,i]
```

Problems?

# Block method

```
for i=myfirstrow..mylastrow  
    x[i] = (y[i] - sum over j=myfirstrow..i-1 ell[i,j]*x[j])  
          / a[i,i]
```

Block Jacobi with local GS solve




Figure : Sparsity pattern corresponding to a block Jacobi preconditioner



# Multicolour ILU

